

Final Analysis Report — Intelligent Hangman (HMM + DQN)

1. Key Observations

- The Bigram HMM oracle with a Greedy agent performed strongly and consistently across a large corpus.
- The DQN agent, trained for 2000 episodes, underperformed the HMM baseline on success rate and wrong guesses.
- Reward shaping and ϵ -greedy exploration helped the DQN learn some useful behavior, but training time and representation likely limited performance.
- Zero repeated guesses in both approaches indicates good masking and action handling.

2. Strategies

- Bigram HMM: Chosen for lightweight, context-aware probabilistic modeling. It captures letter-to-letter dependencies and conditions on the visible pattern, giving strong letter rankings without heavy training.
- DQN (RL): Chosen to learn a policy that can potentially outperform greedy selection by planning better letter orders. It uses an ϵ -decay exploration schedule (start=1.0 \rightarrow end=0.05) to balance exploration and exploitation.
- Reward design: +1 correct, -2 wrong, +10 win, -10 lose, -0.1 per step. This prioritizes accuracy and sample efficiency while penalizing slow solves and mistakes.

3. Exploration vs. Exploitation

- During training, ϵ -greedy encourages trying different letters so the agent can learn Q-values for more actions.
- As ϵ decays, the agent exploits learned Q-values to choose higher-value letters.
- In practice, the DQN needed more training to consistently exploit the oracle-informed state; with only 2000 episodes, it still explored many weak choices.

4. Future Improvements

- Train longer (10k–50k episodes) with periodic evaluations; consider curriculum learning (short words first).
- Improve state features (e.g., positional encodings) and/or use a hybrid policy: restrict actions to top-k oracle letters, select among them via Q-values.
- Try prioritized replay, larger buffer, and tuned rewards (e.g., bonus for early correct streaks).
- Explore trigram language models or subword features to provide richer probabilistic context to both agents.

5. Final Results Comparison Table

Model	Success Rate	Wrong Guesses	Repeated	Final Score
HMM + Greedy	95.0%	3956	0	170,220
DQN Agent	58.45%	9093	0	71,435

Notes:

- Final Score = (SuccessRate \times 2000) – (Wrong \times 5) – (Repeated \times 2)
- Dataset: data/corpus.txt (\approx 50k words). Baseline and DQN evaluated on 2000 games.

6. Train vs. Test Results (with dataset labels)

6.1 Corpus (in-distribution sampling from data/corpus.txt)

- Setup: 2000 games sampled from corpus; lives=6; seed=42.
- Purpose: in-distribution performance ("train/corpus" setting).

Model	Dataset	Games	Success Rate	Wrong Guesses	Repeated	Final Score
HMM + Greedy	corpus.txt sample	2000	95.0%	3956	0	170,220
DQN (Hybrid)	corpus.txt sample	2000	58.45%	9093	0	71,435

Notes:

- DQN uses hybrid evaluation (oracle top \blacktriangleleft k=5) for a fairer comparison and to reduce poor actions.

6.2 Test (held-out list from data/test_words.txt) — OFFICIAL

- Setup: full test_words.txt (n_games = number of lines, here 2000); lives=6; seed=42.
- Purpose: out-of-distribution/held-out performance ("test" setting). These are the official results.

Model	Dataset	Games	Success Rate	Wrong Guesses	Repeated	Final Score
HMM + Greedy	test_words.txt	2000	32.00%	10477	0	11,615
DQN (Strict)	test_words.txt	2000	2.45%	11939	0	-54,795

Notes:

- Baseline is pure argmax (no top \blacktriangleleft k); DQN evaluation is strict (no hybrid restriction).

7. About Score vs. Accuracy (what judges see)

- Final Score = (SuccessRate \times 2000) – (Wrong \times 5) – (Repeated \times 2). It is the primary metric; accuracy (success rate) is a key component but not the only one.
- Why some teams get negative scores: many wrong guesses (large penalty) with low success can outweigh the success reward, turning the score negative.
- Why our corpus score is high but test score is modest/negative: the corpus sample is in-distribution (oracle is strong there). The held-out test is harder; with 6 lives, errors accumulate quickly and dominate the score.